

Whitepaper on Compute

By Anje Kruger, Founder & CEO, Hibarri

“A glimpse into my research on AI & the computational resources required by the future.”

Index

1. Capitalism and Compute.....	1
2. The Basics of the Digital Economy.....	1
4. Economic Analysis of the Bot Economy.....	2
5. Superintelligence.....	3
6. The Function of Bots in Society.....	4
7. What is Compute?.....	5
8. Digital Real Estate.....	6
9. Deficit in the Supply of Compute.....	7
10. Gold Rush to Compute Rush.....	8
11. CPU Cloud Computing.....	10
12. The Cloud Landlord & Tenant Economy in 2030.....	11
13. How Money is Made off Compute.....	11
14. Who Sells Compute.....	12
15. ROI in Real Estate vs Cloud Estate.....	13
16. Cloud Partners.....	14

1. Capitalism and Compute

- 1.1. This thesis practically summarizes compute as a resource, while analyzing what the digital economy currently is, what it is becoming, what it likely looks like in future, how it currently operates, what it currently needs and what it will need in future, for the purpose of identifying the economic opportunity to capitalize on compute as a resource.
- 1.2. The central discovery of this analysis is that the digital economy is a landlord and tenant economy, currently worth \$1.1 trillion and projected to be worth 100 times more in future (literally).
- 1.3. The consequent practical applications of this knowledge include how to capitalize on the digital landlord tenant economy, and how to become a landlord of digital real estate in such an economy.

2. The Basics of the Digital Economy

- 2.1. There is no software and no internet without computers.

- 2.2. The internet is basically a bunch of computers sending signals to each other through networks, like wifi, 4G and 5G.
- 2.3. Computers are basically a processing unit plus a bunch of components that help the processing unit (brain) to do its job. So the processing unit *is* the computer. It computes. Even a cell phone is just a computer with a screen.
- 2.4. The language models we interact with today (ChatGPT, Grok, Claude, Gemini, and Copilot) are advanced computers, or rather, big computers with many subcomputers.
- 2.5. These language models are *not retail products* meant for end consumers to prompt. They were developed to be the *infrastructure of the future internet*.
- 2.6. The future internet is not developers making systems. It is developers or individuals instructing bots on which systems or websites or apps to make. This will be known as the bot economy - a secondary economy of bots performing work digitally for humans.
- 3. The 2030 Bot Economy
 - 3.1. The bot economy will be responsible for scaling the total size of the internet by at least *100 to 1,000 times* in the next 10 years.
 - 3.2. According to Elon Musk, the entire digital economy will become completely driven by agents with almost no further human involvement eventually. Experts estimate that this new machine economy or bot economy will make up more than 90% of the internet in the next 10 years.
 - 3.3. It may very well be that we will see a single individual employing 100 digital employees to build 10 - 20 digital apps, as we are already seeing with GrokBot.
 - 3.4. Individual productivity will increase tenfold across the internet, and so will the compute resource consumption per individual if these agents are adopted at the expected rates.
- 4. Economic Analysis of the Bot Economy
 - 4.1. The “economy” is basically the production of goods and services.
 - 4.2. The current world economy is \$100 trillion large.
 - 4.3. Superintelligence will, no doubt, be an incredibly resource hungry creature. It will require a lot of power, a lot of chips, and a lot of inference technology just to exist as a being.
 - 4.4. With the volume of AI employees and robots increasing, the total amount of working hours increasing, and the productivity per person per hour increasing

tenfold with bots and LLMs, an economist may argue that the world economy will scale by 10 to 100 times to become worth \$1 quadrillion or \$10 quadrillion very soon.

- 4.5. Even if individual digital productivity only increases to become 10 times more per day, that alone will scale the digital economy by $\times 10$.
- 4.6. Fortunately, we are not concerned with the *value* of the economy or the digital economy. We are simply concerned with the volume and value of *resources required for such an economy to function*. Which is currently \$1.1 trillion in resource consumption requirement annually.
- 4.7. Even if the entire economy will still be worth the same in monetary value and never increases in value to become a quadrillion dollar economy despite having increased 10 times in size or digital productivity, it does not concern us because one fact is true; **10 times more digital output per person will require 10 times more input resources per person irrespective of the value of the overall economy.**
- 5. Superintelligence
 - 5.1. Superintelligence certainly holds the spotlight, or 'SI' for short, as opposed to AI.
 - 5.2. Basically, Superintelligence is the state that Artificial Intelligence will officially obtain once it can operate independently and no longer relies on humans, at which point in time its knowledge and capabilities will exceed that of the greatest human mind.
 - 5.3. In order for artificial intelligence to become superintelligence, it would need to be able to become a "doer" - a "bot".
 - 5.4. This is primarily digital superintelligence and digital bots, however, total superintelligence will only be possible once artificial intelligence can move atoms, and pick things up, via their robot bodies (most likely as humanoid robots).
 - 5.5. It sounds very Sci-Fi like, but the reality is that it already exists at a scale of over 100,000 humanoid robots internationally. This scale will likely climb to 100,000,000 within 10 - 20 years from now, in most standard economic models.
 - 5.5.1. Nvidia has already launched software for AI operators to train to become real robots in virtual reality with replica rules and environments that mimic the real world. Once the AI training is complete, the software is downloadable into hundreds of thousands of robot bodies at scale.

- 5.5.2. Tesla is training Optimus with visual inference technology and plans to sell 1 million robots per year from 2027 - 2028.
- 5.5.3. Boston Dynamics are selling their robot dogs at scale already.
- 5.6. At the point where superintelligence exists in both the digital and physical reality (physical reality representing the subset of humanoid robots alone), the resource consumption will scale by another order of magnitude or two, and further, the resource requirements of the international cloud as a collective will become ginormous, mostly because the robots will have to communicate with their operator, computer, or with the cloud to download updates and function in society.
- 5.7. But digital and physical bots serve different purposes and functions.
- 5.8. Calculators are way faster at calculating than even the greatest human mind. Bots, when set up correctly, will become the same, however, for a wider variety of digital tasks.
- 5.9. For the sake of narrowing down this analysis, we will ignore the physical superintelligence related compute requirements that may come of humanoid robots, and focus purely on digital artificial intelligence and bots. In terms of narrowing it down even further, we will look at a world as basic as it was a year ago, with one developer renting cloud computing on one server in the form of cloud compute.
- 6. The Function of Bots in Society
 - 6.1. Eventually, bots will be physical humanoid robots. But for the purpose of this study, when we refer to bots, we refer to digital bots.
 - 6.2. Bots will eventually initiate purchases based on defined parameters, in order to complete the development of systems by buying digital services, the labour of other agents, and compute capacity to execute tasks.
 - 6.3. Very simply, bots will do this because it allows humans to achieve 10 times more per minute or hour than before by doing less of the work.
 - 6.4. There are two very important infrastructural components here. Firstly, the infrastructure that will be required for the language models to run, and secondly, the infrastructure the bots will be using to exist, and to complete tasks.
 - 6.5. Again, we are not concerned with the philosophies of how the bot economy will change society. We are purely analyzing the economic reality of today and tomorrow, and the input requirements of that reality.

- 6.6. Even had we imagined a world without any bots at all, the same infrastructure will continue being required, although at smaller scales.
- 6.7. Bots are basically “doers” powered by “large language models” (LLMs). A “bot” does not think. It basically carries out defined tasks while the “thinking” is done by the underlying large language model, like ChatGPT.
- 7. What is Compute?
 - 7.1. The AI boom has resulted in massive hype about data centers and superintelligence and “compute”. But what is compute?
 - 7.2. Compute is everything that is required for a software system to produce a result, it is also everything required for a language model to “think, produce a thought, infer, make an inference”.
 - 7.3. The purpose of compute is to allow software to process data, basically. It is hardware and electricity, arranged in a specific way so as to be able to produce an output.
 - 7.4. Compute is basically the world’s most consumed commodity.
 - 7.5. Compute matters to every developer because it affects the quality, speed, accuracy and overall functions of their systems. Basically, high quality compute is like a new smart phone and low quality compute is like a 2010 model iPhone.
 - 7.6. Data can be anything - letters, numbers, questions, tasks.
 - 7.7. Language models are basically ¹trillions of data points sitting in memory ²that centralize and organize to produce an answer ³based on memory, principles, and lessons. This “organizing” activity is compute. It is a calculation.
 - 7.8. “Inferences” are currently defined by Google Cloud as “the operational phase where a trained artificial intelligence model processes new, unseen data to generate real-world predictions, decisions, or outputs”.
 - 7.9. While AI training is a valid field and necessary part of the process, the longer term ongoing infrastructure requirement lies in continuous inferences, since the AI will continuously have to “think” - or otherwise “compute”.
 - 7.10. These “inferences” can only be done inside a computer - and a computer is basically a processing unit (usually an NPU, CPU or GPU) or a group of processing units that processes billions of datapoints to “make an inference”.
 - 7.10.1. CPU - Central Processing Unit - processes basically every action
 - 7.10.2. GPU - Graphics Processing Unit - 3D, video, voice, image, complex math
 - 7.10.3. NPU - Neural Processing Unit - a new AI-specific processing unit

- 7.10.4. QPU - Quantum Processing Unit - quantum computing processor
- 7.11. Many investors have started looking at quantum computing as a theoretically “lucrative” investment option, but this has been going since about 2001. The reality is that, quantum computing is mostly in the research and development stages right now, where the cost of a single quantum computer currently exceed \$2.3 million. While it may, no doubt, become a lucrative investment option eventually, too little is known or proven to make definite findings in this market.
- 7.12. Supercomputing is a more mature, measurable and commercially established market. A supercomputer is basically just a bunch of computers computing together. Most LLMs are supercomputers running off interconnected CPUs and GPUs to make inferences.
- 7.13. A computer also has other components, but when developers rent cloud computing, they rent mostly the results produced by the CPU/GPU/NPU and the secondary factor, which is RAM.
- 8. Digital Real Estate
 - 8.1. There is one basic fact to all businesses - they must convert limited resources into products and processes that achieve their goals. For this reason, a company can not do everything. It must do the one thing it is designed to do, and outsource the rest.
 - 8.2. This basic fact of limited resources gave rise to the \$1.1 trillion cloud computing industry a long time ago, shortly after the formation and boom of the internet.
 - 8.3. The economics are relatively straightforward. Cloud computing is financially feasible for software companies as “tenants” or “users” today because it only charges the user for what they actually use on a monthly or hourly basis, so when the computer is not computing, the user does not pay, and there are no upfront fees. It’s all metered.
 - 8.4. The immediate savings can be as enormous as paying \$2.34 per hour for limited compute resources instead of buying a \$212,000 server, because that is all the company needs at that point in time.
 - 8.5. Long term, at maximum capacity, cloud computing works out to be far more expensive, but in the immediate term, it makes sense. And for bots, it is the only available option.
 - 8.6. The cloud computing industry exists because buying servers, setting them up, engineering their rules, handling their cooling, etc, is a gigantic process, not to

mention the major upfront investment requirements. With limited resources, software companies elect to pay per hour to outsource it to cloud service providers and they pay only for what they use, while they focus on developing, selling and improving their products.

- 8.7. Companies can spend years using cloud computing before entering into agreements with data centers for their own server racks, so the question of how long a cloud computing customer keeps renting is usually 3 - 7 years, with their requirements usually increasing by the month.
- 8.8. Data centers, in fact, usually do not sell compute. They sell something called “ping, power and pipe” (rack space, cooling, power and internet) to cloud service providers - who break up, block off and sell the compute. Data centers basically just manage electricity, security and cooling for the servers. Similarly to how landlords rent to co-working spaces, data centers rent out to cloud computing companies.
- 8.9. Historically, cloud computing has just been CPUs, but GPUs are increasingly becoming integrated into this offering, with CPUs going for about \$0.02 per hour for 2GB of RAM, and GPUs going for \$2.34 per hour.
- 8.10. Cloud computing is especially important for software and AI companies because their infrastructure requirements can change dramatically. A company might have 100 customers today and 100,000 customers next month or vice versa. It needs the ability to increase and decrease its computing capacity without having to purchase and install thousands of processors in advance.
- 8.11. Admittedly, there is currently a far narrower pool of AI / language model companies to supply GPU compute to than the general populus of developers looking for CPU based cloud compute. Probably a 1:10,000 ratio. AI cloud customers are extremely high value customers, while there is still a major opportunity for cloud computing as the internet’s total productivity scales up.
- 9. Deficit in the Supply of Compute
 - 9.1. The cloud computing industry currently generates \$1.1 trillion (half as much as the oil & gas industry) in a \$100 trillion world economy.
 - 9.2. With the volume of AI employees and robots increasing, the world economy could scale by 10 to 100 times. The problem is a *major supply shortage* in processing power that already existed before the impending bot economy arose.

- 9.3. Current cloud computing supply cannot keep up with the exploding demand for artificial intelligence and what that means for the world.
- 9.4. Various tech giants have already admitted to having run out of GPUs, CPUs, NPUs, TPUs etc.
- 9.5. GPUs are consistently *sold out* and tech giants like Google, OpenAI and SpaceXAI have already started building their own chips in an attempt to secure more chip supply. Instead of a gold rush, we now have a compute rush.
- 9.6. The opportunity to get a rack slot in data centers is also becoming narrower and narrower as most data centers are already filled to the brim and struggling to get approvals to expand facility sizes to accommodate more rack space.
- 9.7. There are only about 12,000 data centers in the world, while there's an enormous legion of over 50 million software developers and over 1 billion websites on the internet. This does not even take into account the colossal power consumption of large language models and the AI bots that are about to (literally) run the whole internet.
- 9.8. The current world data center supply is already 97 GigaWatts short and another 450 GigaWatts expected in short supply in the coming years. To put that into perspective, a single GigaWatt can power over 1 million homes. 450 GigaWatts can power about 450 million homes. That is the current world deficit.
- 9.9. In that, the \$1.1 trillion cloud computing industry may become \$11 trillion very soon, or \$111 trillion depending on how the value of goods and services keep up with the cost of providing them in the economy.
- 9.10. The opportunity to acquire, own and rent out servers onto the cloud is becoming narrower each day, with established tech networks dominating the acquisition channels and direct investors being completely shut out. But there are methods around this.
- 10. Gold Rush to Compute Rush
 - 10.1. The compute rush has created an enormous opportunity for investors seeking smaller scale opportunities to earn relatively steady, yet scalable capital.
 - 10.2. While there are significant and monumental shifts in the upper echelons amongst the original tech giants, a secondary opportunity has been born for the average investor. It does not involve hefty investments in speculative stocks, hopeful loan agreements, or major risks and it is far more cost effective than real estate.

- 10.3. The average investor can, with the right equipment, partners and network, buy servers to rent them out.
- 10.4. Per Unit Value Analysis
 - 10.4.1. A standard 192 core CPU server sells for about \$13,000 at the moment.
 - 10.4.2. Each of the 192 cores has 2 threads, meaning it has 384 computational “brains”.
 - 10.4.3. With the standard RAM the server comes with, there is usually only 64 GB of RAM available in the small ones. If the resource is packaged as 1 vCPU + 1 GB RAM, physical memory - not CPU threads - would initially limit us to approximately 64 simultaneously guaranteed allocations..
 - 10.4.4. 1GB of RAM (one vCPU) tends to sell for \$0.02 per hour.
 - 10.4.5. $\$0.02 \text{ per vCPU per hour} \times 24 \text{ hours} \times 64 \text{ allocations} \times 30 \text{ days in a month} = \$921 \text{ per month earned on an investment of } \$13,000.00$
 - 10.4.6. But wait, there’s more.
 - 10.4.7. A single user hardly ever exceeds 30% of their total capacity all at once, so sellers overcommit to sell 3 times the total available space, and earning up to \$2,765 per month.
- 10.5. AI Compute Per Unit Analysis
 - 10.5.1. A GPU server sells for far more than a CPU server, and starts within the ranges of \$40,000 per unit, while selling compute on a per-hour basis of \$2.50. The math is relatively similar, only the per-hour yield of a GPU is far superior.
 - 10.5.2. See it as making money every time an AI “thinks”.
 - 10.5.3. And they will be doing a lot of thinking.
- 10.6. AI Inference Economic Analysis
 - 10.6.1. Gartner forecast that worldwide spending on AI-optimised Infrastructure as a Service (IaaS) would reach approximately \$42.3 billion in 2026, representing growth of approximately 96% over 2025.
 - 10.6.2. Gartner also forecast that inference-related infrastructure spending would reach \$23.3 billion in 2026, exceeding spending on AI training at approximately \$19 billion.
 - 10.6.3. The AI inference economy is therefore creating demand for infrastructure that can provide processing capacity continuously, reliably and at a commercially competitive price.

- 10.6.4. It is also important to understand that not every inference requires an expensive GPU. Large language models often use GPUs for computationally intensive inference, but CPUs remain essential for application execution, data processing, orchestration, databases, API requests, agent operations and many smaller AI models.
- 10.6.5. A single AI product can consume several different types of infrastructure. It might use a GPU to generate an answer, a CPU to execute code, additional CPU resources to browse the internet, RAM to retain working data, storage to save information, etc.
- 10.7. The opportunity is not limited to becoming an AI company. There is an equally important opportunity to become one of the infrastructure companies that AI businesses depend on.
- 11. Cloud Computing for Bots
 - 11.1. CPU Cloud Computing is the main infrastructure requirement in the Bot Economy
 - 11.2. While AI is exciting, there won't be that many thinker LLMs, there will mostly be "doer" bots that run off the LLM infrastructure and still 90% of all tasks on the internet will be CPU based.
 - 11.3. Firstly, cloud computing eliminates the need to purchase hardware, and purchasing hardware is significantly more difficult for a bot than for a human, while cloud compute is effortless and easy to select, buy and program for a bot.
 - 11.4. A human developer might operate five applications. A bot might be instructed to create, operate and maintain 100 applications.
 - 11.5. A human might only require a computer for a few hours each day. A digital employee could be instructed to work continuously, including overnight, on weekends and during public holidays.
 - 11.6. A human might manually perform a task once. A bot might perform that task 10,000 times, simultaneously, across multiple applications and customers.
 - 11.7. That creates a massive and almost unimaginable CPU workload.
 - 11.8. The bot does not need to own the physical computer it uses. It simply needs access to the computing capacity required to complete its assigned task.
 - 11.9. This makes cloud computing a natural must-have infrastructure layer for the bot economy.
 - 11.10. With cloud computing, bots can do their jobs and purchase additional resources without a human having to physically install another computer.

12. The Cloud Landlord & Tenant Economy in 2030
 - 12.1. No doubt, the digital landlord tenant economy already exists, but it is about to get far greater than 100 times as big as it currently is!
 - 12.2. The current AI infrastructure conversation is heavily concentrated around GPUs. This is understandable because GPUs are essential to training and running many of the most advanced AI models.
 - 12.3. However, a GPU is generally used for AI inference while a CPU is the general-purpose processing unit that executes.
 - 12.4. LLMs think through GPUs. Bots do tasks through CPUs. Both need RAM, storage and networking.
 - 12.5. Humans currently make up over 95% of cloud tenants and humans make up 100% of cloud landlords. In the next 10 years, humans will still make up 100% of the cloud landlord space, while bots will make up over 50% of the cloud tenant space.
13. Monetizing Compute
 - 13.1. The primary stage money in compute is made by the Data Center who rents out its rack space, cooling, electricity and network to a cloud landlord. We are not interested in this side of the value chain.
 - 13.2. The secondary stage money comes in where the cloud landlord rents to the cloud tenant. RAM and CPU are two of the primary resources sold by cloud-computing providers.
 - 13.3. The capacity the cloud landlord sells could be, as an example, a server with 192 physical CPU cores, 384 hardware threads and approximately 9.2TB (9,200GB) of RAM (more RAM than in the previous calculation). If the cloud landlord chooses to sell its capacity as 384 virtual-compute allocations, with each allocation containing 1 vCPU and 24GB of RAM, the server could theoretically support 384 such allocations without exceeding its physical RAM capacity.
 - 13.4. At a selling price of \$0.02 per 1-vCPU/24GB-RAM allocation per hour, the theoretical maximum gross revenue would be $\$0.02 \times 384 \text{ allocations} \times 24 \text{ hours} \times 30 \text{ days} = \$5,529.60$ per server per month.
 - 13.5. This is a theoretical 100%-occupancy revenue figure, before accounting for CPU utilisation, storage, networking, electricity, data-center costs, software, hardware depreciation, support, payment costs and other operating expenses.

- 13.6. The cloud landlord may also push for CPU overcommitment, since customers generally do not use their maximum CPU allocation continuously. This often happens at a 3 x server capacity level, meaning, instead of making \$5,529.60 per server per month, the cloud landlord pumps out over \$15,000 per server per month.
- 14. Competitors Who Sell Compute
 - 14.1. There are major compute sellers who call themselves “hyperscalers” - which the means is simply that someone can arrive and use only \$5 of compute per month, while having access to \$50,000 of compute the very next day. These major companies (cloud landlords) rent and own data centers, and then rent out their capacity from there. They include (with their 2025 revenue in brackets):
 - 14.1.1. Microsoft Azure (\$75B)
 - 14.1.2. Amazon AWS (\$125B)
 - 14.1.3. Oracle Cloud (\$24B)
 - 14.1.4. Google Cloud Platform (\$58B)
 - 14.1.5. Alibaba Cloud (\$16.8B)
 - 14.1.6. IBM Cloud (Not separately disclosed)
 - 14.1.7. Meta Platforms (Not separately disclosed)
 - 14.1.8. Bytedance (Not separately disclosed)
 - 14.1.9. +10s more
 - 14.2. While the hyperscalers seem like somewhat of an oligopoly, they are actually not. They just have huge customer bases, but their prices are hardly ever competitive and they are often beaten on price by companies such as:
 - 14.2.1. CoreWeave (\$5.1B)
 - 14.2.2. Lambda (\$700M)
 - 14.2.3. Runpod (n.d.)
 - 14.2.4. Fluidstack (n.d.)
 - 14.2.5. Hertzner (n.d.)
 - 14.2.6. +100s more
 - 14.3. Is it still then possible to make your service offering competitive when joining the cloud computing race? Highly.
- 15. ROI in Real Estate vs Cloud Estate
 - 15.1. Real Estate Yields

- 15.1.1. The average rental apartment in New York sells for a minimum of \$900,000 cash and yields \$3,900 per month ROI.
- 15.1.2. An apartment in San Francisco sells for \$800,000 and yields \$3,700 per month.
- 15.1.3. A Texas apartment in Houston costs \$500,000 and yields \$1,500 per month.
- 15.1.4. North London looks a bit better with \$500,000 invested and \$2,750 yields.
- 15.2. Cloud Estate Yields
 - 15.2.1. The average server sells for \$15,000 (small CPU server) and yields \$432 per month in ROI at \$0.02 per vCPU hour with 24GB RAM (this can be done with a Dell AMD EPYC 7502P).
 - 15.2.2. 10 servers at an investment of \$150,000 yields \$4,320 per month on average.
 - 15.2.3. 30 servers at an investment of \$450,000 total yields \$12,960 per month in cash at low performance.
 - 15.2.4. 60 servers at \$900,000 yield \$25,920 per month in cash at low performance.
 - 15.2.5. It is not that the investor has to buy 60 servers. They can also buy 6 servers with 10 times more compute power, and in fact, these will sell at a discount.
 - 15.2.6. At high performance, yields reach \$1,500 per server per month, \$15,000 per month with 10 servers, and \$45,000 per month with 30 servers, and \$90,000 per month on 60 servers.
 - 15.2.7. This assumes one tenant in an apartment in the named areas vs 30 servers selling compute capacity for \$0.012 per 24GB per hour. It also assumes that the rack rental for the server has been paid up front for 24 months in the initial investment. It further assumes that the server yields very low returns and runs at extremely low capacity.
- 15.3. Cloud yields are demonstrably far superior than real estate yields, while requiring far lower investment amounts.
- 15.4. The downsides of real estate are known already, but the main downsides of cloud investment are usually as follows:
 - 15.4.1. Low chance of borrowing or leverage since you can not borrow money for the investment up front without having proven returns. (But, on the plus

side of that coin, there is no interest on the server since you buy it cash, where a property may come at a 10% or 20% interest rate.)

15.4.2. You may not get any rentals on your particular server with the wrong configurations or wrong cloud partner. (This is becoming increasingly easy to avoid with cloud marketplaces and neoclouds).

15.4.3. Rapid depreciation since the server's components rapidly get worn out, with an average server life expectancy of 7 - 10 years at maximum performance. (A \$15,000 server thus loses about \$2,000 of value per year. But on the positive side, if you own a business, the server can be written off fully as a tax write-off.)

16. Cloud Estate

16.1. The question of cloud estate comes down to the same principles as in real estate. Location, timing and the market.

16.2. As for *location*, the best locations to host a server are usually in an economy with high compute demand, such as the USA or UK or India. But it is also an economy closest to the investor where they are able to plug the server out and sell it if needed.

16.3. As for *timing*, the market is currently constrained and demand is huge while supply is scarce and expensive. Buying a server within the next week or month means that it gets rented out quite rapidly, while buying in 12 months from now may mean that you've missed the wave.

16.4. As for the *market*, the market consists of a handful of server suppliers and millions of users, and with current server production capacity being far lower than AI thinking and expansion capacity, suppliers will only keep pushing their prices up in the coming months.

17. Cloud Partners

17.1. Coco by Hibarri allows individual investors to join our cloud partnership program whereby they acquire and own their own servers, and we configure, integrate, and rent those out via our cloud network as a service.

17.2. We help interested investors select and acquire the right servers to serve the growing market of cloud compute users, while buying at the right time before prices go up.

- 17.3. CoCo focuses on integrating deeply in developer networks, marketing our cloud computing network, and driving revenue for server owners in our cloud partner platform. We achieve this by:
 - 17.3.1. Acquiring and setting up servers for our cloud partners
 - 17.3.2. Managing data center relations and billing
 - 17.3.3. Coding instructions, limitations and usage rules into server hardware
 - 17.3.4. Marketing CoCo as a leading cloud compute provider
 - 17.3.5. Making CoCo available to bots through our machine readable interface
 - 17.3.6. Acquiring customers on behalf of our cloud partners
 - 17.3.7. Competitively pricing cloud compute packages
 - 17.3.8. Collecting payments through metering technology
 - 17.3.9. Invoicing on behalf of cloud partners
 - 17.3.10. Metering developer use and adjusting billing per cycle
 - 17.3.11. Distributing workload efficiently across the cloud of servers in our network
 - 17.3.12. Managing all the admin required for efficient payment collections
- 17.4. We charge a 20% margin on the hourly rate billed to the end consumer for continuous management services. This incentivises us to continue marketing and acquiring customers for our network of server owners.
- 17.5. CoCo is an advanced AI NeoCloud and general cloud computing infrastructure for developers, but we are also a marketing partner for private server owners.

- *Anje Kruger*



Resources

1. <https://www.nvidia.com/en-us/use-cases/robotics-simulation/>
2. <https://www.bain.com/insights/ai-data-center-boom-can-we-build-it-if-they-come-technology-report-2026/>
3. <https://datacentrereview.com/2026/04/emea-data-centre-vacancy-hits-record-low-as-ai-demand-outpaces-supply/>
4. <https://www.thenews.com.pk/latest/1415914-elon-musks-spacexai-team-to-build-company-live-using-grok-bot>
5. <https://investingnews.com/the-compute-shortage-threatening-ai-s-next-growth-phase/>
6. <https://spectraip.net/vps/>
7. <https://cloud.google.com/workstations/pricing>
8. <https://www.oracle.com/cloud/pricing/>

Disclaimer: The terms “cloud landlord” and “cloud tenant” are definitions developed by our team for the purposes of this document and are not official or generally recognized industry terminology. Nothing contained in this document constitutes financial, investment, legal, business or other professional advice, nor does it constitute an offer, recommendation or solicitation to invest. This document is provided solely for educational and informational purposes and is intended to present the subject matter from an objective perspective.